

Programming Language Semantics

Gavin Mendel-Gleason

Data Chemist Ltd.

gavin@datachemist.com

Abstract

This paper comprises a set of notes for a short course on programming language semantics using Agda [?]. We cover the basics of operational and denotational semantics, with some example proofs and proof exercises.

Most other proof assistants / programming languages would work as a replacement, and the diligent student should be able to translate them into their language of choice with little difficulty, aside from some subtleties with coinduction in a very limited number of proofs.

2012 ACM Subject Classification Program semantics

Keywords and phrases denotational semantics, operational semantics, type theory

Digital Object Identifier 10.4230/LIPIcs...

1 Introduction

Programming language semantics is the study of the meaning of programs. A programming language can be thought of as some syntactic means of describing what computation we would like to perform, coupled with a mechanism for performing the computation according to what this language describes.

We are going to take a quick tour of programming language semantics, making use of the Agda proof assistant as a method of both implementing, and proving properties of programming languages.

The advantages in clarity of using a formal proof assistant which is itself a programming language as the approach to describing semantics are many. It's often much clearer what precisely is being described in language semantics when we can actually see the datastructures, proofs and evaluation itself described in a well structured metalanguage. We can steal the computation of our host system to clarify meaning as well as give succinct descriptions of what precisely we are trying to prove by means of *types*.

Prior familiarity with type theory is not required to understand this course, though it would be useful if the student is somewhat fluent with both functional programming in a language such as haskell, ML or similar, as well as some familiarity with imperative programming paradigms.

Type theory is a discipline in mathematics which gives a formal approach to the statement of conjectures and evidence for their satisfaction. It is also, however, possible using something known as the Curry-Howard correspondence, to view the evidence of satisfaction, or proofs, as a functional programming language. This is exploited to great effect in so-called 'Dependently typed languages' such as Agda.

We will see how Agda can be used to describe programming language semantics using "Dependent types" for programming languages which we implement in agda.

If all of this is clear as mud: fear not. The detail may in fact dispel the devil.

2 Preliminaries

Instead of thinking of type theory through a very general mathematical lens, we're going to take a more prosaic approach. We will treat it as merely a sophisticated functional



© Gavin Mendel-Gleason;

licensed under Creative Commons License CC-BY

Leibniz International Proceedings in Informatics

LIPICs Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl Publishing, Germany

XX:2 Programming Language Semantics

44 programming language.

45 We start from the ground up, building some basic machinery ourselves in order to
46 understand how it works. Later we will use the analogous machinery from the Agda standard
47 library.

48 Our exploration begins by defining a number of *inductive types*. Inductive types are
49 the primary manner in which we describe our data structures for our semantics. They will
50 also prove useful for defining properties and relations about the terms in the programming
51 languages we will define. Inductive types are essentially abstract data types which are
52 potentially recursive.

53 Inductive types come equipped automatically with mechanisms to introduce elements of
54 them, known as *constructors* and ways of eliminating them, using pattern matching in a
55 manner which should be familiar users of ML and Haskell.

56 The simplest inductive type is the empty type, which we call "bottom" or \perp .

```
57  
58 data  $\perp$  : Set where  
59  
60 explosion :  $\forall \{A : \text{Set}\} \rightarrow \perp \rightarrow A$   
61 explosion ()  
62  
63  $\neg\_ :$   $\forall (A : \text{Set}) \rightarrow \text{Set}$   
64  $\neg A = A \rightarrow \perp$   
65
```

66 This type has very little to say. You can't produce one so it has no constructors, and if
67 you find one, you can't pattern match on it because it could not have been introduced in the
68 first place.

69 We can see how this is used with our term `explosion` which is the proof of the *principle of*
70 *explosion*. First, the brackets, `{ }` around the `A` say that the argument is *implicit* and we are
71 asking Agda to try and infer it from context. This means that the first argument we pattern
72 match on is the terms of type \perp . However, since there are none to match, we don't even
73 have to finish our description of what to do if we find an occurrence and so need not find a
74 way to produce an element of the type `A`. From false premises, anything follows.

75 In addition we can describe a kind of negation with `¬_` which allows us to say that our
76 premise can not hold, because otherwise we would be able to produce a datatype with no
77 elements. This will prove a powerful, if somewhat non-intuitive way of describe that certain
78 properties do not hold.

79 The next simplest type is the type called *top* or \top . Sometimes this type is called *one*
80 (and the type \perp is called *zero*) since it only has the one constructor: `tt`.

```
81  
82 data  $\top$  : Set where  
83   tt :  $\top$   
84
```

85 Next we see the more familiar type of booleans. The boolean type has two constructors,
86 `true` and `false`.

```
87 data Bool : Set where  
88   true : Bool  
89   false : Bool
```

```

90
91   not : Bool → Bool
92   not true = false
93   not false = true
94
95   _and_ : Bool → Bool → Bool
96   true and true = true
97   true and false = false
98   false and b = false
99
100  _or_ : Bool → Bool → Bool
101  true or b = true
102  false or true = true
103  false or false = false
104
105  if_then_else_ : ∀ {C : Set} → Bool → C → C → C
106  if true then x else y = x
107  if false then x else y = y
108

```

With the `bool` type defined, we can define the `not` function, which maps booleans to booleans, but flips the constructor.

Similarly we can define the usual `_and_` and `_or_` which allow us to combine booleans in the usual way. Here the underscores tell Agda that we mean for the function to be infix, so its first argument is on the left, and its second is on the right.

However, we can also describe mixfix functions in Agda, as we see with the `if_then_else_` function, which takes a boolean as its first argument, and then takes two branches of type `C` which the boolean is used to select from. We can then recover the quite familiar form of conditional branching as a function.

For a more in-depth description of the Agda programming language of dependent types and its use as a proof-assistant, see Ulf Norell's work *Dependently Typed Programming in Agda*[?] and Aaron Stump's *Verified Functional Programming in Agda*[?]. We will press on hoping that the reader can pick things up as we go along, and, if not, looks to these additional resources in the event of confusion.

3 Operational Semantics

Now that we've seen some basic uses of inductive types and functions which manipulate them, we are ready to do some heavier lifting in the service of programming language semantics.

operational semantics is an approach to semantics which looks at how, operationally, our programming terms *compute*. we'll get into the nitty-gritty of evaluation, and learn how to state general properties of terms by seeing how evaluation progresses.

```

129
130   module operationalsemantics where
131     open import Data.Nat
132
133     data Exp : Set where
134       num : ℕ → Exp
135       _⊕_ : Exp → Exp → Exp
136

```

We define a new *module* which we name `OperationalSemantics`. In Agda, a module is essentially a record that allows us to define parameters, control our namespace and produce a collection of associated terms which refer to each other. In this case we have a very simple module with no parameters. Directly below the introduction of our module, and within its scope, we import another module, `Data.Nat` which contains the definitions for the natural numbers, and various functions to manipulate them and facts about them.

We can then define the syntax of our first language: `Exp`. This language is composed of numbers, and addition. The syntax for introducing a number uses the constructor `num` on a natural number, the definition of which we've imported from `Data.Nat`. We can then add two expressions using the infix constructor `_⊕_`.

So what does an expression look like? The expression `twelve` shows how to write the addition of 3 numbers in our syntax.

```

149
150 twelve : Exp
151 twelve = num 3 ⊕ (num 4 ⊕ num 5)
152

```

Our syntax is, however, merely a data-structure. It doesn't *do* anything. We must add *dynamics* to our syntax in order to get something worthy of being called a language.

We can do this by defining what it would mean to *evaluate* an expression. We will do this by defining an *evaluation relation*.

3.1 Big-Step Operational Semantics

The first evaluation relation we will define, is something called a *big-step evaluation relation*. It's called a big-step relation because we define a relation between a term in our syntax, with the final stage of evaluation is when something has reached a *value*. In this case our values are numbers.

This provides a contrast with *small-step evaluation relations* that instead tell you what the next step of a computation is. We will look at small-step evaluation relations later, and then explore how they relate.

```

infix 10 _⇓_
data _⇓_ : Exp → ℕ → Set where
  n⇓n : ∀ {n} →

-----

num n ⇓ n

E⊕E : ∀ {E1 E2 n1 n2} →
  E1 ⇓ n1 → E2 ⇓ n2 →
  -----
  E1 ⊕ E2 ⇓ (n1 + n2)

```

166 We introduce a new, infix, inductive type which we name `_↓_`. We will call this a relation,
 167 because it has two parameters, an expression, drawn from `Exp` and a natural number drawn
 168 from `ℕ`.

169 All other types we've seen were defined to be themselves of type `Set` directly. We can think
 170 of this as a family of types, with each member of the family drawn by choosing elements of
 171 `Exp` and `ℕ`, or as is natural for our current problem, as a *relation* between `Exp` and `ℕ`.

172 Our relation has two constructors. The first is called `n↓n`. It basically says, that the
 173 big step evaluation relation relates an expression formed of a natural number, to that same
 174 natural number. In other words, when you evaluate a natural number, nothing happens.

175 We use a number of dashes, `--` to simulate the horizontal bar sometimes used in proofs,
 176 to separate our premises, from our conclusions. In Agda more than two dashes is simply a
 177 comment, so this is just a visual aid and has no impact on the terms we are expressing.

178 In our first case, we have no premises (aside from the uninteresting implicit `n`).

179 The second constructor is `E⊕E`. We can read this as stating that, given we know the
 180 number to which two expressions E_1 and E_2 evaluate, we can then determine that the
 181 expressions conjoined with `_⊕_` will evaluate to the sum of the numbers to which each
 182 expression evaluates, respectively.

```
183 evalBig : ∀ E → Σ[ n ∈ ℕ ] E ↓ n
evalBig (num x) = x , n↓n
evalBig (e ⊕ e₁) with evalBig e | evalBig e₁
evalBig (e ⊕ e₁) | n , proof_n | m , proof_m = n + m , E⊕E proof_n proof_m
```

184 With this bigstep evaluation relation in hand, we can write an evaluation function, `evalBig`,
 185 which demonstrates that given any expression in our syntax defined above, we can obtain a
 186 natural number which relates the expression to its final evaluated form.

187 In Agda, we do this using the `Σ` record. This is a pair datastructure, comprising an
 188 element, and a proof of the type $E ↓ n$ which depends on both n and E . In this case the
 189 element is some natural number `ℕ` and the proof we must construct is the full demonstration
 190 that E evaluates to that number under the big step evaluation relations.

191 In English we might read the type of `evalBig` as stating that:

192 “Given any expression, `evalBig` will give us the number this expression evaluates to,
 193 and a proof that this is a value which is related by the big-step evaluation relation.”

194 So how does `evalBig` work? It is essentially a simple proof by induction on the structure
 195 of the syntax. With the syntax of our language we have two cases. If we are the base case,
 196 namely `(num x)` we can use the fact that numbers evaluate to themselves. This expression
 197 has no sub-expressions, and so we are done, which is what makes it a base case.

198 If we are a sum of two expressions, combined with the `_⊕_` constructor, we first evaluate
 199 each sub expression independently, and then sum their values, using the `E⊕E` constructor to
 200 combine the proofs obtained in both branches to produce the *next step* of the proof.

201 With `evalBig` in hand, we can look back at `twelve` and see the effect of evaluating this
 202 expression.

```

203   example↓ : num 3 ⊕ (num 4 ⊕ num 5) ↓ 12
        example↓ = proj₂ (evalBig (num 3 ⊕ (num 4 ⊕ num 5)))

```

204 We apply `evalBig` to our expression, and then use `proj₂` to project out the *proof* from our
 205 pair of the number 12 and the proof which it is coupled with.

206 This example provides a sort of template for the way in which we can build up operational
 207 semantics using big-step evaluation. Using this general approach as a guide, we can look at
 208 more interesting languages, and ask more sophisticated questions and get proofs about our
 209 semantics.

210 However, for now we will continue on with this simple language, and look at another
 211 approach to operational semantics.

212 3.2 Small-step operational semantics

213 In our previous example we demonstrated our reduction relation by relating expressions to
 214 the values which they must evaluate to. There is another option, namely, that we describe
 215 our reduction in terms of *the next step*.

216 We can define a new type, `_→_`, which relates two expressions when the expression on
 217 the left can transition to the expression on the right in *one step*.

```

infix 8 _→_
data _→_ : Exp → Exp → Set where
  +→ : ∀ {n m} →

-----
num n ⊕ num m → num (n + m)

⊕₁→ : ∀ {E₁ E₁' E₂} →
  E₁ → E₁' →
-----
  E₁ ⊕ E₂ → E₁' ⊕ E₂

⊕₂→ : ∀ {n E₂ E₂'} →
  E₂ → E₂' →
-----
  num n ⊕ E₂ → num n ⊕ E₂'

```

219 This time we have three rules with one base case and a left and right rule. The base case,
 220 `+→` takes two values combined with `_⊕_` and evaluates to a value by summing the two
 221 values.

222 The left rule, `⊕₁→`, relates an expression with a sum in which the left-hand summand
 223 can be further related by a small step, but whose right-hand summand remains unchanged.

224 The right rule, `⊕₂→`, relates an expression with a sum in which the left-hand is already

225 fully evaluated to a number, but whose right-hand can still be related by a small step.

226 To get a feel for whats going on, we can look at how we can relate expressions concretely
227 using this relation.

228
$$\begin{aligned} \text{example} \longrightarrow_1 & : (\text{num } 3 \oplus \text{num } 7) \oplus (\text{num } 8 \oplus \text{num } 1) \longrightarrow \text{num } 10 \oplus (\text{num } 8 \oplus \text{num } 1) \\ \text{example} \longrightarrow_1 & = \oplus_1 \longrightarrow + \longrightarrow \\ \text{example} \longrightarrow_2 & : (\text{num } 10) \oplus (\text{num } 8 \oplus \text{num } 1) \longrightarrow \text{num } 10 \oplus \text{num } 9 \\ \text{example} \longrightarrow_2 & = \oplus_2 \longrightarrow + \longrightarrow \end{aligned}$$

229 In $\text{example} \longrightarrow_1$ the left hand side has summed exactly one pair of numbers. This is done
230 by descending into the left branch and then summing the pair of summands.

231 In $\text{example} \longrightarrow_2$ we have already reduced the left-hand side to a number, we can then use
232 the $\oplus_2 \longrightarrow$ to descend into the right branch where we find a pair of numbers which can be
233 summed.

234 The reader may have noticed that these rules are written very carefully such that there is
235 no choice in the application of rules. You have to evaluate all left expressions first, before
236 one is allowed to proceed with the right branch.

237 This choice is not necessary, but describes a *deterministic* relation. Instead of this we
238 could choose an alternative approach which allows a choice of summand in which to descend.

infix 8 $\longrightarrow_{\text{ch}}$
data $\longrightarrow_{\text{ch}}$: $\text{Exp} \rightarrow \text{Exp} \rightarrow \text{Set}$ where
 $+ \longrightarrow_{\text{ch}}$: $\forall \{n\ m\} \rightarrow$

$$\text{num } n \oplus \text{num } m \longrightarrow_{\text{ch}} \text{num } (n + m)$$

$\oplus_1 \longrightarrow_{\text{ch}}$: $\forall \{E_1\ E_1'\ E_2\} \rightarrow$

$$E_1 \longrightarrow_{\text{ch}} E_1' \rightarrow$$

$$E_1 \oplus E_2 \longrightarrow_{\text{ch}} E_1' \oplus E_2$$

$\oplus_2 \longrightarrow_{\text{ch}}$: $\forall \{E_1\ E_2\ E_2'\} \rightarrow$

$$E_2 \longrightarrow_{\text{ch}} E_2' \rightarrow$$

$$E_1 \oplus E_2 \longrightarrow_{\text{ch}} E_1 \oplus E_2'$$

240 We have a new relation, $\longrightarrow_{\text{ch}}$, where the *ch* is short for 'choice'. Here we have the
241 same approach to taking two summands which are themselves values and combining it to a
242 single value. However for the left and right rules, we are not constrained in which branch we
243 choose. We can choose to make our step in either the right branch or the left branch.

XX:8 Programming Language Semantics

```

example→ch1 : (num 3 ⊕ num 7) ⊕ (num 8 ⊕ num 1) →ch num 10 ⊕ (num 8 ⊕ num 1)
example→ch1 = ⊕1→ch +→ch
244 example→ch2 : (num 3 ⊕ num 7) ⊕ (num 8 ⊕ num 1) →ch (num 3 ⊕ num 7) ⊕ num 9
example→ch2 = ⊕2→ch +→ch

```

245 We can now see two different examples where we make a choice of branch in which to step.
 246 In `example→ch1` we chose the left branch for evaluation, summing 3 and 7 and producing
 247 10. In `example→ch2` we take the right branch, summing 8 and 1 to produce 9.

248 Now that we have defined these two different approaches, we might ask, are we performing
 249 the same evaluation?

```

→⇒→ch : ∀ {E1 E2} → (E1 → E2) → (E1 →ch E2)
→⇒→ch +→ = +→ch
250 →⇒→ch (⊕1→ d) = ⊕1→ch (→⇒→ch d)
→⇒→ch (⊕2→ d) = ⊕2→ch (→⇒→ch d)

```

251 We can see from the theorem `→⇒→ch` that we can always mimic the deterministic
 252 small step evaluation relation `→` with the `→ch` by simply always choosing the
 253 semetric proof rule.

254 However, the reverse theorem which transforms a choice small step to the deterministic
 255 one is not true. To prove this, we make use of Agda's `Data.Empty` and the `Relation.Nullary`
 256 libraries, rather than our own negation which we described in the introduction.

```

open import Data.Empty
open import Relation.Nullary

257 ¬→ch⇒→ : ¬ (∀ E1 E2 → (E1 →ch E2) → (E1 → E2))
¬→ch⇒→ f with f ((num 0 ⊕ num 0) ⊕ (num 0 ⊕ num 0)) ((num 0 ⊕ num 0) ⊕ num 0)
              (⊕2→ch +→ch)
¬→ch⇒→ f | ()

```

258 To prove that we can not always perform the same small step reduction we need merely
 259 to find a single example for which does not hold.

260 Since `¬_` is essentially short hand for a function which maps from a type to the empty
 261 type, our first argument, `f` has the type of the theorem

```

262 ∀ E1 E2 → (E1 →ch E2) → (E1 → E2)

```

263 ... which we know can not be true.

264 With our proof in hand, we can apply it to a patently impossible reduction. Agda
 265 cleverly detects that we have no cases in which this can hold and we are able to eliminate the
 266 possibility of our argument, which eliminates the need to supply an impossible proof of `⊥`.

267 We have claimed earlier that our small step evaluation relation is deterministic, but we
 268 have not yet proved it. In order to prove determinism, it is convenient first to have a notion
 269 of equivalence. In Agda we can do this by means of a propositional equality type `≡p`.

270 We will define our own here in order to get a flavour. Later we will use Agda's libraries to
 271 handle this machinery for us.

```

272 infix 4 _≡_
data _≡_ {A : Set} (x : A) : A → Set where
instance reflp : x ≡p x

```

273 The equality type can be described as follows. Given any type and a value of that type,
 274 we can show that the value is equivalent to itself, using the `reflp` constructor. This might seem
 275 quite useless, as we can only show things are equal to themselves! However, it is more flexible
 276 than it first appears, because Agda is able to perform computations. So in order to show
 277 two things are equal, we will be able to make the host system perform some computation
 278 and we can build up non-trivial equivalences in this way.

279 Now that we have the basic idea of equivalence, we can switch to Agda's built in
 280 propositional equality `_≡_` and its single constructor `refl`.

```

open import Relation.Binary.PropositionalEquality

→deterministic : ∀ {E E1 E2} → E → E1 → E → E2 → E1 ≡ E2
→deterministic +→ +→ = refl
→deterministic +→ (⊕1 → ())
→deterministic +→ (⊕2 → ())
281 →deterministic (⊕1 → ()) +→
→deterministic (⊕1 → p) (⊕1 → q) = cong2 _⊕_ (→deterministic p q) refl
→deterministic (⊕1 → ()) (⊕2 → q)
→deterministic (⊕2 → ()) +→
→deterministic (⊕2 → p) (⊕1 → ())
→deterministic (⊕2 → p) (⊕2 → q) = cong2 _⊕_ refl ((→deterministic p q))

```

282 That our small step evaluation relation is deterministic is described by the type of
 283 `→deterministic`. The theorem we are trying to prove, is that given any term E in our
 284 language, and a proof that this is related by a single step to a term E_1 , then if it is also
 285 related by a single step to some other term E_2 then E_1 and E_2 must actually be the same
 286 term.

287 The proof of the theorem proceeds by induction on the proof of both reduction relations
 288 using a simultaneous case match. The base case of both is trivial, giving us that the sum of
 289 the numbers n and m which are implicit variables of `+→` are equal to themselves which is
 290 easily proved with `refl`. Though we can not see these numbers in the proof as written here,
 291 replacing `refl` with a *hole* by compiling Agda with a `'?` will allow you to see the implicit
 292 variables in the context.

293 Six more of the cases are eliminated by Agda automatically and can be replaced with
 294 a vacuous match. As an example, in the first case in which we have a vacuous match, the
 295 argument to `⊕1 →` would have to have type: `num n → E1`. Case analysis on this yields no
 296 ways in which to form this type as `num n` can not be the left-hand side of any reduction.

297 In order to deal with the recursive case, where we have like constructors, we can make use
 298 of `conf2`. This function merely uses the fact that if two arguments are equal, then functions

XX:10 Programming Language Semantics

299 of those arguments are also equal.

300