

# Programming Language Semantics

Gavin Mendel-Gleason

Data Chemist Ltd.

mendelgg@scss.tcd.ie

## Abstract

This paper comprises a set of notes for a short course on programming language semantics using Agda [?]. We cover the basics of operational and denotational semantics, with some example proofs and proof exercises.

Most other proof assistants / programming languages would work as a replacement, and the diligent student should be able to translate them into their language of choice with little difficulty, aside from some subtleties with coinduction in a very limited number of proofs.

**2012 ACM Subject Classification** Program semantics

**Keywords and phrases** denotational semantics, operational semantics, type theory

**Digital Object Identifier** 10.4230/LIPIcs...

## 1 Introduction

Programming language semantics is the study of the meaning of programming languages. A programming language can be thought of as some syntactic means of describing what computation we would like to perform, coupled with a mechanism for performing the computation according to what this language describes.

We are going to take a quick tour of programming language semantics, making use of the Agda proof assistant as a method of both implementing, and proving properties of programming languages.

The advantages in clarity of using a formal proof assistant which is itself a programming language as the approach to describing semantics are many. It's often much clearer what precisely is being described in language semantics when we can actually see the datastructures, proofs and evaluation itself described in a well structured metalanguage. We can steal the computation of our host system to clarify meaning as well as give succinct descriptions of what precisely we are trying to prove by means of *types*.

Prior familiarity with type theory is not required to understand this course, though it would be useful if the student is somewhat fluent with both functional programming in a language such as haskell, ML or similar, as well as some familiarity with imperative programming paradigms.

Type theory is a discipline in mathematics which gives a formal approach to the statement of conjectures and evidence for their satisfaction. It is also, however, possible using something known as the Curry-Howard correspondence, to view the evidence of satisfaction, or proofs, as a functional programming language. This is exploited to great effect in so-called 'Dependently typed languages' such as Agda.

We will see how Agda can be used to describe programming language semantics using "Dependent types" for programming languages which we implement in agda.

If all of this is clear as mud: fear not. The detail may in fact dispel the devil.

## 2 Preliminaries

Instead of thinking of type theory through a very general mathematical lens, we're going to take a more prosaic approach. We are



© Gavin Mendel-Gleason;

licensed under Creative Commons License CC-BY

Leibniz International Proceedings in Informatics

LIPICs Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl Publishing, Germany

## XX:2 Programming Language Semantics

44 We'll start from the ground up, building some basic machinery ourselves in order to  
45 understand how it works. Later we will use the analogous machinery from the Agda standard  
46 library.

47 We will start our exploration by defining a number of *inductive types*. Inductive types are  
48 the primary manner in which we describe our data structures for our semantics. They will  
49 also prove useful for defining properties and relations about the terms in the programming  
50 languages we will define.

51 Inductive types come equipped automatically with mechanisms to introduce elements of  
52 them, known as *constructors* and ways of eliminating them, using pattern matching in a  
53 manner which should be familiar users of ML and Haskell.

54 The simplest inductive type is the empty type, which we call "bottom" or  $\perp$ .

```
55  
56 data  $\perp$  : Set where  
57  
58 explosion :  $\forall \{A : \text{Set}\} \rightarrow \perp \rightarrow A$   
59 explosion ()  
60  
61  $\neg\_ :$   $\forall (A : \text{Set}) \rightarrow \text{Set}$   
62  $\neg A = A \rightarrow \perp$   
63
```

64 This type has very little to say. You can't produce one so it has no constructors, and if  
65 you find one, you can't pattern match on it because it could not have been introduced in the  
66 first place.

67 We can see how this is used with our term `explosion` which is the proof of the *principle of*  
68 *explosion*. First, the brackets, `{ }` around the `A` say that the argument is *implicit* and we are  
69 asking Agda to try and infer it from context. This means that the first argument we pattern  
70 match on is the terms of type  $\perp$ . However, since there are none to match, we don't even  
71 have to finish our description of what to do if we find an occurrence and so need not find a  
72 way to produce an element of the type `A`. From false premises, anything follows.

73 In addition we can describe a kind of negation with `¬_` which allows us to say that our  
74 premise can not hold, because otherwise we would be able to produce a datatype with no  
75 elements. This will prove a powerful, if somewhat non-intuitive way of describe that certain  
76 properties do not hold.

77 The next simplest type is the type called *top* or  $\top$ . Sometimes this type is called *one*  
78 (and the type  $\perp$  is called *zero*) since it only has the one constructor: `tt`.

```
79  
80 data  $\top$  : Set where  
81   tt :  $\top$   
82
```

83 Next we see the more familiar type of booleans. The boolean type has two constructors,  
84 `true` and `false`.

```
85 data Bool : Set where  
86   true : Bool  
87   false : Bool  
88  
89   not : Bool  $\rightarrow$  Bool
```

```

90  not true = false
91  not false = true
92
93  _and_ : Bool → Bool → Bool
94  true and true = true
95  true and false = false
96  false and b = false
97
98  _or_ : Bool → Bool → Bool
99  true or b = true
100 false or true = true
101 false or false = false
102
103 if_then_else_ : ∀ {C : Set} → Bool → C → C → C
104 if true then x else y = x
105 if false then x else y = y
106

```

With the `bool` type defined, we can define the `not` function, which maps booleans to booleans, but flips the constructor.

Similarly we can define the usual `_and_` and `_or_` which allow us to combine booleans in the usual way. Here the underscores tell Agda that we mean for the function to be infix, so its first argument is on the left, and its second is on the right.

However, we can also describe mixfix functions in Agda, as we see with the `if_then_else_` function, which takes a boolean as its first argument, and then takes two branches of type `C` which the boolean is used to select from. We can then recover the quite familiar form of conditional branching as a function.

For a more in-depth description of the Agda programming language of dependent types and its use as a proof-assistant, see Ulf Norell's work *Dependently Typed Programming in Agda*[?] and Aaron Stump's *Verified Functional Programming in Agda*[?]. We will press on hoping that the reader can pick things up as we go along, and, if not, looks to these additional resources in the event of confusion.

### 3 Operational Semantics

Now that we've seen some basic uses of inductive types and functions which manipulate them, we are ready to do some heavier lifting in the service of programming language semantics.

operational semantics is an approach to semantics which looks at how, operationally, our programming terms *compute*. we'll get into the nitty-gritty of evaluation, and learn how to state geeneral properties of terms by seeing how evaluation progresses.

```

127
128 module operationalsemantics where
129   open import Data.Nat
130
131   data Exp : Set where
132     num : ℕ → Exp
133     _⊕_ : Exp → Exp → Exp
134

```

## XX:4 Programming Language Semantics

135 We define a new *module* which we name `OperationalSemantics`. In Agda, a module is  
136 essentially a record that allows us to define parameters, control our namespace and produce  
137 a collection of associated terms which refer to eachother. In this case we have a very simple  
138 module with no parameters. Directly below the introduction of our module, and within its  
139 scope, we import another module, `Data.Nat` which contains the definitions for the natural  
140 numbers, and various functions to manipulate them and facts about them.

141 We can then define the syntax of our first language: `Exp`. This language is composed of  
142 numbers, and addition. The syntax for introducing a number uses the constructor `num` on a  
143 natural number, the definition of which we've imported from `Data.Nat`. We can then add  
144 two expressions using the infix constructor `_⊕_`.

145 So what does an expression look like? The expression `twelve` shows how to write the  
146 addition of 3 numbers in our syntax.

```
147  
148 twelve : Exp  
149 twelve = num 3 ⊕ (num 4 ⊕ num 5)  
150
```

151 Our syntax is, however, merely a data-structure. It doesn't *do* anything. We must add  
152 *dynamics* to our syntax in order to get something worthy of being called a language.

153 We can do this by defining what it would mean to *evaluate* an expression. We will do  
154 this by defining an *evaluation relation*.

### 155 3.1 Big-Step Operational Semantics

156 The first evaluation relation we will define, is something called a *big-step evaluation relation*.  
157 It's called a big-step relation because we define a relation between a term in our syntax, with  
158 the final stage of evaluation is when something has reached a *value*. In this case our values  
159 are numbers.

160 This provides a contrast with *small-step evaluation relations* that instead tell you what  
161 the next step of a computation is. We will look at small-step evaluation relations later, and  
162 then explore how they relate.

```
infix 10 _⇓_  
data _⇓_ : Exp → ℕ → Set where  
  n⇓n : ∀ {n} →  
  
-----  
  num n ⇓ n  
163  
  
E⇓E : ∀ {E1 E2 n1 n2} →  
  
  E1 ⇓ n1 → E2 ⇓ n2 →  
-----  
  E1 ⊕ E2 ⇓ (n1 + n2)
```

164 We introduce a new, infix, inductive type which we name `_⇓_`. We will call this a relation,  
165 because it has two parameters, an expression, drawn from `Exp` and a natural number drawn  
166 from `ℕ`.

167 All other types we've seen were defined to be themselves of type `Set` directly. We can think  
 168 of this as a family of types, with each member of the family drawn by choosing elements of  
 169 `Exp` and `ℕ`, or as is natural for our current problem, as a *relation* between `Exp` and `ℕ`.

170 Our relation has two constructors. The first is called `n↘n`. It basically says, that the  
 171 big step evaluation relation relates an expression formed of a natural number, to that same  
 172 natural number. In other words, when you evaluate a natural number, nothing happens.

173 We use a number of dashes, `--` to simulate the horizontal bar sometimes used in proofs,  
 174 to separate our premises, from our conclusions. In our first case, we have no premises (aside  
 175 from the uninteresting implicit `n`).

176 The second constructor is `E⊕E`. We can read this as stating that, given we know the  
 177 number to which two expressions  $E_1$  and  $E_2$  evaluate, we can then determine that the  
 178 expressions conjoined with `_⊕_` will evaluate to the sum of the numbers to which each  
 179 expression evaluates, respectively.

```
180 evalBig : ∀ E → Σ[ n ∈ ℕ ] E ↘ n
    evalBig (num x) = x , n↘n
    evalBig (e ⊕ e₁) with evalBig e | evalBig e₁
    evalBig (e ⊕ e₁) | n , proof_n | m , proof_m = n + m , E⊕E proof_n proof_m
```

181 With this bigstep evaluation relation in hand, we can write an evaluation function, `evalBig`,  
 182 which demonstrates that given any expression in our syntax defined above, we can obtain a  
 183 natural number which relates the expression to its final evaluated form.

184 In Agda, we do this using the `Σ` record. This is a pair datastructure, comprising an  
 185 element, and a proof of the type  $E ↘ n$  which depends on both `n` and `E`. In this case the  
 186 element is some natural number `ℕ` and the proof we must construct is the full demonstration  
 187 that `E` evaluates to that number under the big step evaluation relations.

188 In English we might read the type of `evalBig` as stating that:

189 “Given any expression, `evalBig` will give us the number this expression evaluates to,  
 190 and a proof that this is a value which is related by the big-step evaluation relation.”

191 So how does `evalBig` work? It is essentially a simple proof by induction on the structure  
 192 of the syntax. With the syntax of our language we have two cases. If we are the base case,  
 193 namely `(num x)` we can use the fact that numbers evaluate to themselves. This expression  
 194 has no sub-expressions, and so we are done, which is what makes it a base case.

195 If we are a sum of two expressions, combined with the `_⊕_` constructor, we first evaluate  
 196 each sub expression independently, and then sum their values, using the `E⊕E` constructor to  
 197 combine the proofs obtained in both branches to produce the *next step* of the proof.

198 With `evalBig` in hand, we can look back at `twelve` and see the effect of evaluating this  
 199 expression.

```
200 example↘ : num 3 ⊕ (num 4 ⊕ num 5) ↘ 12
    example↘ = proj₂ (evalBig (num 3 ⊕ (num 4 ⊕ num 5)))
```

201 We apply `evalBig` to our expression, and then use `proj₂` to project out the *proof* from our  
 202 pair of the number `12` and the proof which it is coupled with.

203 This example provides a sort of template for the way in which we can build up operational  
 204 semantics using big-step evaluation. Using this general approach as a guide, we can look at  
 205 more interesting languages, and ask more sophisticated questions and get proofs about our  
 206 semantics.

207 However, for now we will continue on with this simple language, and look at another  
 208 approach to operational semantics.

## 209 3.2 Small-step operational semantics

210 In our previous example we demonstrated our reduction relation by relating expressions to  
 211 the values which they must evaluate to. There is another option, namely, that we describe  
 212 our reduction in terms of *the next step*.

213 We can define a new type,  $\_ \longrightarrow \_$ , which relates two expressions when the expression on  
 214 the left can transition to the expression on the right in *one step*.

```

infix 8  $\_ \longrightarrow \_$ 
data  $\_ \longrightarrow \_ : \text{Exp} \rightarrow \text{Exp} \rightarrow \text{Set}$  where
   $+ \longrightarrow : \forall \{n\ m\} \rightarrow$ 

$$\frac{}{\text{num } n \oplus \text{num } m \longrightarrow \text{num } (n + m)}$$

 $\oplus_1 \longrightarrow : \forall \{E_1\ E_1'\ E_2\} \rightarrow$ 

$$\frac{E_1 \longrightarrow E_1'}{E_1 \oplus E_2 \longrightarrow E_1' \oplus E_2}$$

 $\oplus_2 \longrightarrow : \forall \{n\ E_2\ E_2'\} \rightarrow$ 

$$\frac{E_2 \longrightarrow E_2'}{\text{num } n \oplus E_2 \longrightarrow \text{num } n \oplus E_2'}$$


```

216 This time we have three rules with one base case and a left and right rule. The base case,  
 217  $+ \longrightarrow$  takes two values combined with  $\_ \oplus \_$  and evaluates to a value by summing the two  
 218 values.

219 The left rule,  $\oplus_1 \longrightarrow$ , relates an expression with a sum in which the left-hand summand  
 220 can be further related by a small step, but whose right-hand summand remains unchanged.

221 The right rule,  $\oplus_2 \longrightarrow$ , relates an expression with a sum in which the left-hand is already  
 222 fully evaluated to a number, but whose right-hand can still be related by a small step.

223 To get a feel for what's going on, we can look at how we can relate expressions concretely  
 224 using this relation.

225 `example→1 : (num 3 ⊕ num 7) ⊕ (num 8 ⊕ num 1) → num 10 ⊕ (num 8 ⊕ num 1)`  
`example→1 = ⊕1→ +→`  
`example→2 : (num 10) ⊕ (num 8 ⊕ num 1) → num 10 ⊕ num 9`  
`example→2 = ⊕2→ +→`

226 In `example→1` the left hand side has summed exactly one pair of numbers. This is done  
 227 by descending into the left branch and then summing the pair of summands.

228 In `example→2` we have already reduced the left-hand side to a number, we can then use  
 229 the `⊕2→` to descend into the right branch where we find a pair of numbers which can be  
 230 summed.

231 The reader may have noticed that these rules are written very carefully such that there is  
 232 no choice in the application of rules. You have to evaluate all left expressions first, before  
 233 one is allowed to proceed with the right branch.

234 This choice is not necessary, but describes a *deterministic* relation. Instead of this we  
 235 could choose an alternative approach which allows a choice of summand in which to descend.

236 `infix 8 _→ch_`  
`data _→ch_ : Exp → Exp → Set where`  
`+→ch : ∀ {n m} →`  

$$\text{num } n \oplus \text{num } m \rightarrow_{\text{ch}} \text{num } (n + m)$$
  

$$\oplus_1 \rightarrow_{\text{ch}} : \forall \{E_1 \ E_1' \ E_2\} \rightarrow$$
  

$$E_1 \rightarrow_{\text{ch}} E_1' \rightarrow$$
  

$$E_1 \oplus E_2 \rightarrow_{\text{ch}} E_1' \oplus E_2$$
  

$$\oplus_2 \rightarrow_{\text{ch}} : \forall \{E_1 \ E_2 \ E_2'\} \rightarrow$$
  

$$E_2 \rightarrow_{\text{ch}} E_2' \rightarrow$$
  

$$E_1 \oplus E_2 \rightarrow_{\text{ch}} E_1 \oplus E_2'$$

237 We have a new relation, `_→ch_`, where the *ch* is short for 'choice'. Here we have the  
 238 same approach to taking two summands which are themselves values and combining it to a  
 239 single value. However for the left and right rules, we are not constrained in which branch we  
 240 choose. We can choose to make our step in either the right branch or the left branch.

## XX:8 Programming Language Semantics

```

example→ch1 : (num 3 ⊕ num 7) ⊕ (num 8 ⊕ num 1) →ch num 10 ⊕ (num 8 ⊕ num 1)
example→ch1 = ⊕1→ch +→ch
241 example→ch2 : (num 3 ⊕ num 7) ⊕ (num 8 ⊕ num 1) →ch (num 3 ⊕ num 7) ⊕ num 9
example→ch2 = ⊕2→ch +→ch

```

242 We can now see two different examples where we make a choice of branch in which to step.  
 243 In `example→ch1` we chose the left branch for evaluation, summing 3 and 7 and producing  
 244 10. In `example→ch2` we take the right branch, summing 8 and 1 to produce 9.

245 Now that we have defined these two different approaches, we might ask, are we performing  
 246 the same evaluation?

```

→⇒→ch : ∀ {E1 E2} → (E1 → E2) → (E1 →ch E2)
→⇒→ch +→ = +→ch
247 →⇒→ch (⊕1→ d) = ⊕1→ch (→⇒→ch d)
→⇒→ch (⊕2→ d) = ⊕2→ch (→⇒→ch d)

```

248 We can see from the theorem `→⇒→ch` that we can always mimic the deterministic  
 249 small step evaluation relation `→` with the `→ch` by simply always choosing the  
 250 semetric proof rule.

251 However, the reverse theorem which transforms a choice small step to the deterministic  
 252 one is not true. To prove this, we make use of Agda's `Data.Empty` and the `Relation.Nullary`  
 253 libraries, rather than our own negation which we described in the introduction.

```

open import Data.Empty
open import Relation.Nullary

254 ¬→ch⇒→ : ¬ (∀ E1 E2 → (E1 →ch E2) → (E1 → E2))
¬→ch⇒→ f with f ((num 0 ⊕ num 0) ⊕ (num 0 ⊕ num 0)) ((num 0 ⊕ num 0) ⊕ num 0)
               (⊕2→ch +→ch)
¬→ch⇒→ f | ()

```

255 To prove that we can not always perform the same small step reduction we need merely  
 256 to find a single example for which does not hold.

257 Since `¬_` is essentially short hand for a function which maps from a type to the empty  
 258 type, our first argument, `f` has the type of the theorem

```

259 ∀ E1 E2 → (E1 →ch E2) → (E1 → E2)

```

260 ... which we know can not be true.

261 With our proof in hand, we can apply it to a patently impossible reduction. Agda  
 262 cleverly detects that we have no cases in which this can hold and we are able to eliminate the  
 263 possibility of our argument, which eliminates the need to supply an impossible proof of `⊥`.

264 We have claimed earlier that our small step evaluation relation is deterministic, but we  
 265 have not yet proved it. In order to prove determinism, it is convenient first to have a notion  
 266 of equivalence. In Agda we can do this by means of a propositional equality type `≡p`.



267 We will define our own here in order to get a flavour. Later we will use Agda's libraries to  
 268 handle this machinery for us.

```

269 infix 4 _≡_p_
data _≡_p_ {A : Set} (x : A) : A → Set where
instance refl_p : x ≡_p x

```

270 The equality type can be described as follows. Given any type and a value of that type,  
 271 we can show that the value is equivalent to itself, using the `refl_p` constructor. This might seem  
 272 quite useless, as we can only show things are equal to themselves! However, it is more flexible  
 273 than it first appears, because Agda is able to perform computations. So in order to show  
 274 two things are equal, we will be able to make the host system perform some computation  
 275 and we can build up non-trivial equivalences in this way.

276 Now that we have the basic idea of equivalence, we can switch to Agda's built in  
 277 propositional equality `_≡_` and its single constructor `refl`.

```

open import Relation.Binary.PropositionalEquality

→deterministic : ∀ {E E1 E2} → E → E1 → E → E2 → E1 ≡ E2
→deterministic +→ +→ = refl
→deterministic +→ (⊕1 → ())
→deterministic +→ (⊕2 → ())
278 →deterministic (⊕1 → ()) +→
→deterministic (⊕1 → p) (⊕1 → q) = cong2 _⊕_ (→deterministic p q) refl
→deterministic (⊕1 → ()) (⊕2 → q)
→deterministic (⊕2 → ()) +→
→deterministic (⊕2 → p) (⊕1 → ())
→deterministic (⊕2 → p) (⊕2 → q) = cong2 _⊕_ refl ((→deterministic p q))

```

279 That our small step evaluation relation is deterministic is described by the type of  
 280 `→deterministic`. The theorem we are trying to prove, is that given any term  $E$  in our  
 281 language, and a proof that this is related by a single step to a term  $E_1$ , then if it is also  
 282 related by a single step to some other term  $E_2$  then  $E_1$  and  $E_2$  must actually be the same  
 283 term.

284 The proof of the theorem proceeds by induction on the proof of both reduction relations  
 285 using a simultaneous case match. The base case of both is trivial, giving us that the sum of  
 286 the numbers  $n$  and  $m$  which are implicit variables of `+→` are equal to themselves which is  
 287 easily proved with `refl`. Though we can not see these numbers in the proof as written here,  
 288 replacing `refl` with a *hole* by compiling Agda with a `'?` will allow you to see the implicit  
 289 variables in the context.

290 Six more of the cases are eliminated by Agda automatically and can be replaced with  
 291 a vacuous match. As an example, in the first case in which we have a vacuous match, the  
 292 argument to `⊕1 →` would have to have type: `num n → E1`. Case analysis on this yields no  
 293 ways in which to form this type as `num n` can not be the left-hand side of any reduction.

294 as they require two different constructors to be equivalent. Since inductive datatypes  
 295 assume that constructors are disjoint by construction, Agda can safely eliminate cases in

## XX:10 Programming Language Semantics

296 which two constructors actually being the same.

297 This leaves us with only

298